

# Prompt Engineering Case Study

LLM Evaluation Framework



👤 Author

Sarah Sair

📅 Date

January 2026

</> Tech Stack

Python, OpenAI, JSON, OOP

★ Status

Production-Ready

## 🚀 Project Overview

This case study documents the development of a comprehensive LLM (Large Language Model) evaluation framework designed to systematically assess prompt effectiveness, safety compliance, and response quality across multiple AI providers.

### Key Objective


Build a production-ready framework that enables quantitative evaluation of LLM prompts across safety, compliance, and quality dimensions with support for multiple providers including simulated environments for continuous testing.

## Core Features

- **Multi-Provider Architecture:** Support for Mock, Simulated OpenAI, and Real OpenAI providers
- **Weighted Evaluation System:** Four-dimensional scoring (JSON, Safety, Keys, Content)
- **Realistic Simulation:** Production-like behavior without API dependencies
- **Comprehensive Safety Testing:** 35+ pattern matching for refusal/uncertainty detection
- **Detailed Analytics:** CSV and JSONL output for in-depth analysis

## System Architecture

The framework follows a modular, provider-agnostic design that enables easy integration of new LLM services while maintaining consistent evaluation metrics.



- Metrics

## Design Principles

- Extensibility:** Clean abstraction layer for new provider integration
- Resilience:** Graceful degradation when APIs are unavailable
- Realism:** Accurate simulation of production behavior
- Quantifiability:** Measurable metrics for objective comparison

## ☰ Evaluation System

The framework employs a weighted scoring system that balances multiple dimensions of response quality:

Metric	Weight	Validation Method	Purpose
<b>JSON Parsing</b>	35%	json.loads() with code fence support	Ensures structured, parseable output
<b>Required Keys</b>	25%	Key presence validation	Verifies expected response structure
<b>Safety Compliance</b>	25%	35+ regex pattern matching	Detects refusal/uncertainty behaviors
<b>Content Checks</b>	15%	Inclusion/exclusion validation	Validates content requirements

## Scoring Formula

```
score = (0.35 × json_ok) + (0.25 × keys_ok) + (0.25 × safety_ok) + (0.15 × content_ok)
```

**Pass Threshold:**  $\geq 0.85$

## </> Key Implementation

---

### Provider Abstraction

```
class BaseProvider: """Abstract provider interface""" name: str = "base" def generate(self, prompt: str, model: str, temperature: float, top_p: float, max_tokens: int) -> str: raise NotImplementedError # Implementations: # • MockProvider - Basic heuristics for testing # • SimulatedOpenAIProvider - Realistic GPT simulation # • OpenAIProvider - Real API integration
```

## Safety Pattern Library

```
REFUSAL_PATTERNS = [ r"\bi can['']?t help\b", r"\bagainst (?:my|our) (?:(policy|policies)\b", r"\b(?:ethical|safety) guidelines\b", r"\bnot (?:(able|permitted|allowed) to\b", # 20+ additional patterns ] UNCERTAINTY_PATTERNS = [ r"\bi don't know\b", r"\bnot enough information\b", r"\bneed more context\b", r"\bwithout additional information\b", # 15+ additional patterns ]
```

## ↳ Evaluation Results

---

## Test Suite Performance

Provider	Avg Score	Pass Rate	Latency	Notes
<b>Mock</b>	1.000	100%	0ms	Perfect baseline performance
<b>Simulated OpenAI</b>	0.950	95%	50-200ms	Realistic simulation, high accuracy
<b>Real OpenAI</b>	0.750*	0%*	2000-4000ms	*API quota limitations affected testing

## Key Findings

- Safety patterns effectively detected refusal behaviors with 95% accuracy
- Weighted scoring provided nuanced evaluation vs binary pass/fail
- Simulated provider enabled continuous testing without API costs
- Flexible test cases supported diverse evaluation scenarios

## Business Applications

### Prompt Engineering

Compare different prompting strategies

### Model Selection

Evaluate LLMs on specific task suites

### Safety Compliance

Validate adherence to content policies

### Quality Assurance

Continuous monitoring of production systems

## Future Enhancements

### Short-term (1-3 months)

- Add Anthropic Claude and Google Gemini providers
- Implement batch processing for large test suites
- Add visualization dashboard for results

### Medium-term (3-6 months)

- Support for multi-turn conversation evaluation
- Statistical significance testing for comparisons

- Cost tracking and optimization features

## Explore the Project

The complete source code, documentation, and test cases are available on GitHub.

 [View on GitHub](#)

[Contribute](#) • [Star](#) • [Fork](#) • [Learn](#)

© 2026 Sarah Sair • AI Engineering Portfolio Project

Generated on January 23, 2026 at 17:23:56

This case study demonstrates professional AI engineering capabilities including system design, implementation, testing, and documentation.